

TOPIC 5

Transaction Management

Upon completion this topic, students should be able to:

5.1 Demonstrate database transaction management

CLO 1 : Apply fundamental of Database Management System (DBMS), relational data model and normalization concepts in database development process.

CLO 2 : Show a well-structured database using the database query to manipulate a database with an appropriate commercial Database Management System.

What is a Transaction?

- A logical unit of work on a database
 - An entire program
 - A portion of a program
 - A single command
- The entire series of steps necessary to accomplish a logical unit of work
- Successful transactions change the database from one CONSISTENT STATE to another
(One where all data integrity constraints are satisfied)

Example of a Transaction

- Updating a Record
 - Locate the Record on Disk
 - Bring record into Buffer
 - Update Data in the Buffer
 - Writing Data Back to Disk

Transaction processing systems category

- a. Batch transaction processing system
 - is an efficient way of processing large volumes of data. where a group of transactions is collected over a period of time. Data is collected, entered, processed and then the batch results are produced.
 - e.g payroll system
- b. On-line transaction processing system (OLTP)
 - is a type of data processing that consists of executing a number of transactions occurring concurrently.
 - e.g online banking, shopping, order entry, or sending text messages.
- c. Real time transaction processing system
 - requires quick transaction and characterized by supplying immediate response.
 - In real time processing processor needs to very responsive and active all the time.
 - e.g ATM transaction, radar system

4 Properties of a Transaction

- Atomic – All or Nothing

All parts of the transaction must be completed and committed or it must be aborted and rolled back

- Consistent

Each user is responsible to ensure that their transaction (if executed by itself) would leave the database in a consistent state

4 Properties of a Transaction

- Isolation

The final effects of multiple simultaneous transactions must be the same as if they were executed one right after the other

- Durability

If a transaction has been committed, the DBMS must ensure that its effects are permanently recorded in the database (even if the system crashes)

Transaction Management with SQL

- SQL Statements → Commit / Rollback
- When a transaction sequence is initiated it must continue through all succeeding SQL statements until:
 1. A Commit Statement is Reached
 2. A Rollback Statement is Reached
 3. The End of the Program is Reached (Commit)
 4. The Program is Abnormally Terminated (Rollback)

Example

```
BEGIN TRAN
    DECLARE @ErrorCode INT, @TranSuccessful INT
    SET @TranSuccessful = 1

    INSERT INTO tblCatalog (CatalogYear)
        VALUES('2002')
    SET @ErrorCode = @@ERROR; IF (@ErrorCode <> 0) SET @TranSuccessful = 0 –
    False

    INSERT INTO tblCatalog (CatalogYear)
        VALUES('2003')
    SET @ErrorCode = @@ERROR; IF (@ErrorCode <> 0) SET @TranSuccessful = 0 –
    False

    IF @TranSuccessful = 0
        BEGIN
            ROLLBACK TRAN
            RAISERROR ('Rolledback transaction: Insert Catalog Year.', 16,1)
        END
    ELSE
        BEGIN
            COMMIT TRAN
            PRINT 'Successfully inserted catalog years...'
        END
GO
```

Transaction Log

- Keeps track of all transactions that update the database
 - Record for the beginning of the transaction
 - Type of operation (insert / update / delete)
 - Names of objects/tables affected by the transaction
 - Before and After Values for Updated Fields
 - Pointers to Previous and Next Transaction Log Entries for the same transaction
 - The Ending of the Transaction (Commit)
- Used for recovery in case of a Rollback

Concurrency Control

- Coordination of simultaneous transaction execution in a multiprocessing database system
- Ensure transaction serializability in a multi-user database
- Lack of Concurrency Control can create data integrity and consistency problems:
 - Lost Updates
 - Uncommitted Data
 - Inconsistent Retrievals

Lost Updates

<u>Time</u>	<u>Jack's Trans</u>	<u>Jill's Trans</u>	<u>Balance</u>
T1	Begin		
T2	Read Balance	Begin	1000
T3		Read Balance	1000
T4	Bal = Bal - 50 (950)		1000
T5	Write Bal (950)	Bal = Bal + 100 (1100)	950
T6	Commit		950
T7		Write Bal (1100)	1100
T8		Commit	1100

Uncommitted Data

Time	Deposit	Interest	Bal
T1	Begin Transaction		1000
T2	Read Bal (1000)		1000
T3	Bal = Bal + 1000 (2000)		1000
T4	Write Bal (2000)	Begin Transaction	2000
T5		Read Bal (2000)	2000
T6		Bal = Bal*1.05 (2100)	2000
T7	Rollback		1000
T8		Write Bal (2100)	2100
T9		Commit	2100

Inconsistent Retrievals

Time	SumBal	Transfer	Bal A	Bal B	Bal C	Sum
T1	Begin Trans		5000	5000	5000	
T2	Sum = 0	Begin Trans	5000	5000	5000	
T3	Read BalA (5000)		5000	5000	5000	
T4	Sum = Sum + BalA (5000)	Read BalA (5000)	5000	5000	5000	
T5	Read BalB (5000)	BalA = BalA -1000 (4000)	5000	5000	5000	
T6	Sum = Sum+BalB (10000)	Write BalA (4000)	4000	5000	5000	
T7		Read BalC	4000	5000	5000	
T8		BalC =BalC + 1000 (6000)	4000	5000	5000	
T9		Write BalC (6000)	4000	5000	6000	
T10	Read BalC	Commit	4000	5000	6000	
T11	Sum=Sum + BalC (16000)		4000	5000	6000	
T12	Write Sum (16000)		4000	5000	6000	16000
T13	Commit		4000	5000	6000	16000

Serial Execution of Transactions

- Serial Execution of transaction means that the transactions are performed one after another.
- No interaction between transactions - No Concurrency Control Problems
- Serial Execution will never leave the database in an inconsistent state → Every Serial Execution is considered correct (Even if a different order would cause different results)

Serializability

- If 2 Transactions are only reading data items –
They do not conflict → Order is unimportant
- If 2 Transactions operate (Read/Write) on
Separate Data Items
 - They do not conflict → Order is unimportant
- If 1 Transaction Writes to a Data Item and
Another Reads or Writes to the Same Data Item
→ The Order of Execution IS Important

The Scheduler

- Special DBMS Program to establish the order of operations in which concurrent transactions are executes
- Interleaves the execution of database operations to ensure:
 - Serializability
 - Isolation of Transactions

The Scheduler

- Bases its actions on Concurrency Control Algorithms (Locking / Time Stamping)
- Ensures the CPU is used efficiently (Scheduling Methods)
- Facilitates Data Isolation → Ensure that 2 transactions do not update the same data at the same time

Concurrency Control Algorithms

- Locking

A Transaction “locks” a database object to prevent another object from modifying the object

- Time-Stamping

Assign a global unique time stamp to each transaction

- Optimistic

Assumption that most database operations do not conflict

Locking

- Lock guarantees exclusive use of data item to current transaction
- Prevents reading Inconsistent Data
- Lock Manager is responsible for assigning and policing the locks used by the transaction

Locking Granularity

Indicates the level of lock use

- Database Level – Entire Database is Locked
- Table Level – Entire Table is Locked
- Page Level – Locks an Entire Diskpage
(Most Frequently Used)
- Row Level – Locks Single Row of Table
- Field Level – Locks a Single Attribute of a
Single Row (Rarely Done)

Types of Locks:

Binary

- Binary Locks – Lock with 2 States
 - Locked – No other transaction can use that object
 - Unlocked – Any transaction can lock and use object

All Transactions require a Lock and Unlock Operation for Each Object Accessed (Handled by DBMS)

- Eliminates Lost Updates
- Too Restrictive to Yield Optimal Concurrency Conditions

Types of Locks:

Shared / Exclusive Locks

- Indicates the Nature of the Lock
- Shared Lock – Concurrent Transactions are granted READ access on the basis of a common lock
- Exclusive Lock – Access is reserved for the transaction that locked the object
- 3 States: Unlocked, Shared (Read), Exclusive (Write)
- More Efficient Data Access Solution
- More Overhead for Lock Manager
 - Type of lock needed must be known
 - 3 Operations:
 - Read_Lock – Check to see the type of lock
 - Write_Lock – Issue a Lock
 - Unlock – Release a Lock
 - Allow Upgrading / Downgrading of Locks

Problems with Locking

- Transaction Schedule May Not be Serializable
 - Can be solved with 2-Phase Locking
- May Cause Deadlocks
 - A deadlock is caused when 2 transactions wait for each other to unlock data

Two Phase Locking

- Defines how transactions Acquire and Relinquish Locks
 1. Growing Phase – The transaction acquires all locks (doesn't unlock any data)
 2. Shrinking Phase – The transaction releases locks (doesn't lock any additional data)
- Transactions acquire all locks it needs until it reaches locked point
- When locked, data is modified and locks are released

Deadlocks

- Occur when 2 transactions exist in the following mode:

T1 = access data item X and Y

T2 = Access data items Y and X

If T1 does not unlock Y, T2 cannot begin

If T2 does not unlock X, T1 cannot continue

T1 & T2 wait indefinitely for each other to unlock data

- Deadlocks are only possible if a transactions wants an Exclusive Lock (No Deadlocks on Shared Locks)

Controlling Deadlocks

- Prevention – A transaction requesting a new lock is aborted if there is the possibility of a deadlock – Transaction is rolled back, Locks are released, Transaction is rescheduled
- Detection – Periodically test the database for deadlocks. If a deadlock is found, abort / rollback one of the transactions
- Avoidance – Requires a transaction to obtain all locks needed before it can execute – requires locks to be obtained in succession

Optimistic Method

- Most database operations do not conflict
- No locking or time stamping
- Transactions execute until commit
 - Read Phase – Read database, execute computations, make local updates (temporary update file)
 - Validate Phase – Transaction is validated to ensure changes will not effect integrity of database
 - If Validated → Go to Write Phase
 - If Not Validated → Restart Transaction and discard initial changes
 - Write Phase – Commit Changes to database
- Good for Read / Query Databases (Few Updates)

Database Recovery

- Restore a database from a given state to a previous consistent state
- Atomic Transaction Property (All or None)
- Backup Levels:
 - Full Backup
 - Differential Backup
 - Transaction Log Backup
- Database / System Failures:
 - Software (O.S., DBMS, Application Programs, Viruses)
 - Hardware (Memory Chips, Disk Crashes, Bad Sectors)
 - Programming Exemption (Application Program rollbacks)
 - Transaction (Aborting transactions due to deadlock detection)
 - External (Fire, Flood, etc)

Transaction Recovery

- Recover Database by using data in the Transaction Log
- Write-Ahead-Log – Transaction logs need to be written before any database data is updated
- Redundant Transaction Logs – Several copies of log on different devices
- Database Buffers – Buffers are used to increase processing time on updates instead of accessing data on disk
- Database Checkpoints – Process of writing all updated buffers to disk → While this is taking place, all other requests are not executes
 - Scheduled several times per hour
 - Checkpoints are registered in the transaction log

Database Backup

- Backups are useful primarily for two purposes. The first is to restore a state following a disaster (called disaster recovery).
- The second is to restore small numbers of files after they have been accidentally deleted or corrupted.
- Data loss is also very common. 66% of internet users have suffered from serious data loss.